

Parallel and Scientific Computing

Coursework Two

Donald Whyte (sc10dw@leeds.ac.uk)

Friday 29th November, 2013

Contents

1	Performance as Problem Size Grows	2
2	Larger Convergence Tolerance and its Effects	2
3	Scalability of Row Decomposition	2
4	Scalability of Block Decomposition	4
4.a	More Scalable Partition Strategy	7
4.b	Relation to Theoretical Parallel Time Complexity Analysis	7
5	Jacobi and Gauss-Seidel Iteration	7
5.a	Lexicographic Gauss-Seidel on a Single Processor	7
5.b	Lexicographic Gauss-Seidel on Multiple Processors	10
6	Red-Black Gauss-Seidel Iteration	11
7	Successive Over-Relaxation	11
8	Comparison of the Three Iterative Methods	12
8.a	Comparison on a Single Processor	12
8.b	Comparison on Multiple Processors	12
8.c	Shortest Runtime Achieved	13
9	Appendix	13
9.a	Task 1 Code Modifications	13
9.b	Task 3 Code Modifications	13
9.c	Task 4 MPD Ring	14
9.d	Task 5 Code Modifications	14
9.e	Task 6 Full Source Code – redblack_gauss.c	14

Problem Size (N)	Error
8	2.59991e+01
16	6.90990e+00
32	1.87038e+00
64	4.76563e-01
128	1.19343e-01
256	2.98689e-02
512	7.46826e-03

Table 1: Contains the error of the final solution produce by `heat_rows.c` for different approximation intervals and a convergence tolerance of 10^{-12}

1 Performance as Problem Size Grows

The exact solution to the problem has been changed to $e^{3y}\cos 3x$, with a convergence tolerance of $10^{-12} = 0.000000000001$. The Laplace solver with Jacobi iteration was run on this problem using various problem sizes. The error of the approximated solutions for each of these problem sizes is given in Table 1. Figure 1 shows a logarithmic plot of the problem size against the magnitude of the error. Notice how this plot is a straight line, going downwards as N is increases.

This matches the values in Table 1. Increasing N decreases the magnitude of the error. Every time the problem size is **doubled** (interval is **halved**), the error is roughly **quartered**.

This means the error is proportional to the interval size. That is, $Error \in O(dx^2)$ where $Error$ is the error of the final solution and dx is the interval size ("spacial" step) being used.

2 Larger Convergence Tolerance and its Effects

The same experiences were performed as the ones in section 1, except with a convergence tolerance of $10^{-5} = 0.00001$. Table 2 and Figure 2 show the errors with different problem/interval sizes and a logarithmic plot of problem size against error respectively.

Initially, the behaviour of the error appears similar to when a convergence tolerance of 10^{-5} is used. Halving the interval size by doubling N both roughly quarters the error repeatedly, until $N = 128$ is reached.

At that point, when 128 is doubled to 256, the error slightly increases. Doubling again to 512 makes the error four times as large. This increase in error after 128 is illustrated in Figure 2. This shows that as you reduce the interval size, the error in the found solution starts to increase after a certain point. Exactly when the error starts increasing depends on what convergence tolerance is used.

The reason for this is larger problem sizes mean smaller interval sizes. This typically increases the accuracy of the found solution, but it means that **smaller changes per iteration** are inevitable. If the convergence tolerance is quite large, then these small changes will cause the algorithm to think it has converged to a solution and terminate. At that point, the current iteration is used as the solution, which may have a large error because it has stopped prematurely.

3 Scalability of Row Decomposition

`heat_rows.c` was modified such that it always runs for exactly 8000 iterations. The execution times of Jacobi iteration with row decomposition for different problem sizes and number of processes were recorded. These times are shown in Table 3.

For each problem size, the execution time to perform the 8000 iterations decreases as you increase the number of processes being used, **until** a certain number of processes is reached. At that number, adding any more processors **increases execution time**.




Figure 1: Logarithmic plot illustrating how error changes as approximation as problem size increases

Problem Size (N)	Error
8	2.59991e+01
16	6.90983e+00
32	1.87012e+00
64	4.75122e-01
128	1.13469e-01
256	1.29022e-01
512	4.44284e-01

Table 2: Contains the error of the final solution produce by heat_rows.c for different approximation intervals and a convergence tolerance of 10^{-5}




Figure 2: Logarithmic plot illustrating how error changes as approximation problem size increases

I suspect the reason for this is that any potential speedup achieved by parallelising the solver is **dominated** by the increased communication costs introduced when adding more processors.

This is supported by the fact that the number of processors it takes to make the solver run slower increases as the problem size increases. This is because the speedup gained by parallelising for large values of N dominates the time it takes to perform the extra communication. With small values of N , the execution time is reduced less, meaning the time spent with inter-process communication (which is the *same* as it is for large values N) is larger than the reduction in time. Therefore, the final execution time for small problem sizes will actually be *longer* than solving the problem with a single processor.

The reduction is also illustrated in Figure 3 and 4, which plot number of processors against the solver's **speedup** and **efficiency** factors respectively (for varying problem sizes). Notice how the speedup achieved by parallelising the solver drops off for each problem size, with the speedup of smaller problem sizes dropping off earlier with less processors.

The efficiency also shows this property; the overall efficiency of this implementation of parallel Jacobi steadily reduces as you increase the number of processors, with larger problem sizes dropping at much slower rates on average.

NOTE: There are some fluctuations in the results (e.g. speedup and efficiency of $N = 961$). This is most likely due to spikes in use of the machines in the MPD ring. That said, the speedup and efficiency factor plots still show the overall relation between problem size and number of processes, so it shouldn't cause too much of an issue.

4 Scalability of Block Decomposition

heat_blocks.c was modified in the same way as **heat_rows.c** in section 3. The execution times of Jacobi iteration with block decomposition for different problem sizes and number of processes were

Problem Size (N)	# Processors	Execution Time (seconds)
61	1	1.51095e-01
61	2	1.06213e-01
61	4	2.52394e-01
61	8	5.70459e+00
61	16	9.74602e+00
121	1	7.06480e-01
121	2	4.10898e-01
121	4	3.53546e-01
121	8	5.35226e+00
121	16	1.07147e+01
241	1	2.50098e+00
241	2	1.46303e+00
241	4	1.03135e+00
241	8	7.08373e+00
241	16	1.52967e+01
481	1	1.51405e+01
481	2	8.39778e+00
481	4	4.23518e+00
481	8	8.64512e+00
481	16	1.64684e+01
961	1	5.34672e+01
961	2	3.43274e+01
961	4	2.46849e+01
961	8	1.53521e+01
961	16	2.22045e+01
1921	1	1.90737e+02
1921	2	1.35852e+02
1921	4	9.96819e+01
1921	8	5.21617e+01
1921	16	3.42827e+01

Table 3: Table illustrating how the problem size and number of processes used affects runtime when using row decomposition




Figure 3: Speedup achieved by parallelising the Laplace solver using row (strip) decomposition for different problem sizes




Figure 4: Efficiency of parallel Laplace solver using row (strip) decomposition for different problem sizes

recorded and placed in Table 4. Speedup and efficiency factor plots are shown in Figures 5 and 6.

Like strip decomposition, the execution time to perform the 8000 iterations decreases as you increase the number of processes being used **until** a certain number of processes is reached. At that number, adding any more processors **increases execution time**.

4.a More Scalable Partition Strategy

So which partition strategy scales better as the number of processors is increased? Strip or block decomposition? Comparing the runtimes directly from Tables 3 and 4, block decomposition looks to be the better solution, as its execution times are consistently smaller than strip decomposition's, regardless of problem size.

However, if we compare the speedup and efficiency factors of the two strategies strip decomposition looks as if it would scale better as the number of processors is increased. The speedup gained by adding more processors is reduced slower with strip decomposition than block decomposition. In fact, Figures 3 and 5 show that the speedup with strip decomposition is still increasing when $N = 1921$ with 16 processors, whereas it has already started to decrease with block decomposition.

The same applies to the efficiency. The efficiency factor appears to start decreasing earlier with strip decomposition, but the rate at which efficiency is reduced is much slower. With $N = 1921$ and 16 processors, $E(p) \approx 0.35$ for strip decomposition and $E(p) \approx 0.19$ for block decomposition.

Therefore, I would conclude that strip decomposition provides better scalability for the Laplace solver.

4.b Relation to Theoretical Parallel Time Complexity Analysis

This contradicts the theoretical parallel time complexity analysis described in lectures. In that, block decomposition appeared to provide better scalability, based on the speedup factor plots shown. However, plots based on real runtime information appear to suggest that strip decomposition provides better overall scalability/efficiency.

In the theoretical analysis, the communications overhead with block decomposition may have been greatly underestimated in the theoretical analysis, or perhaps that the actual algorithm simply has a faster implementation for row decomposition. Processing, memory access and communication times all have an impact on performance, and they all vary based on the platform, the operating system and other processes running on the same machines. That is, the environment the experiments were run in could have been a factor in the produced plots too.

In practice, there is always a point at which adding more processes introduces so much communications overhead such that it dominates any speedup gained by parallelising the problem (as discussed in section 3).

5 Jacobi and Gauss-Seidel Iteration

5.a Lexicographic Gauss-Seidel on a Single Processor

Lexicographic Gauss-Seidel iteration was implemented in the file **lexicographic_gauss.c**. The number of iterations it took for algorithm to converge to the exact solution where $u = 1$ and $N = 241$, with different convergence tolerances of 10^{-12} , were recorded.

These iteration counts are available in Table 5, along with how many iterations to took the original Jacobi solver to be implemented.

As the convergence tolerance is increased, the number of iterations reached to converge to a solution is reduced for **both** Jacobi and Gauss-Seidel. When using the same convergence tolerance, Gauss-Seidel roughly takes **half** the number of iterations to converge to a solution as Jacobi. For example, when $Tol = 10^{-12}$, Jacobi takes 220567 iterations whereas Gauss-Seidel takes $114365 \approx \frac{220567}{2}$ iterations. This is well-known property of Gauss-Seidel and generally holds true for most computational problems.

Problem Size (N)	# Processors	Execution Time (seconds)
61	1	1.08981e-01
61	4	7.82480e-02
61	9	6.11493e+00
61	16	6.73178e+00
121	1	3.38402e-01
121	4	1.77012e-01
121	9	7.14676e+00
121	16	8.37292e+00
241	1	1.21609e+00
241	4	4.80048e-01
241	9	7.11296e+00
241	16	1.07243e+01
481	1	4.50174e+00
481	4	1.48817e+00
481	9	1.11971e+01
481	16	1.39045e+01
961	1	2.08103e+01
961	4	1.20978e+01
961	9	1.65290e+01
961	16	1.93283e+01
1921	1	9.15544e+01
1921	4	5.14857e+01
1921	9	3.01908e+01
1921	16	3.20344e+01

Table 4: Table illustrating how the problem size and number of processes used affects runtime when using block decomposition




Figure 5: Speedup achieved by parallelising the Laplace solver using block decomposition for different problem sizes




Figure 6: Efficiency of parallel Laplace solver using block decomposition for different problem sizes

Algorithm	Convergence Tolerance	# Iterations
Jacobi	10^{-12}	220567
Gauss-Seidel	10^{-12}	114365
Jacobi	10^{-10}	166368
Gauss-Seidel	10^{-10}	87265
Jacobi	10^{-8}	112169
Gauss-Seidel	10^{-8}	60165
Jacobi	10^{-6}	57969
Gauss-Seidel	10^{-6}	33065

Table 5: Shows how many iterations are required by the Jacobi and Gauss-Seidel algorithms to converge to a solution with a particular error tolerance

Algorithm	# Processors	# Iterations
Jacobi	2	220567
Jacobi	3	220567
Jacobi	4	220567
Gauss-Seidel	2	116198
Gauss-Seidel	3	117053
Gauss-Seidel	4	117994

Table 6: Shows how many iterations are required by the Jacobi and Gauss-Seidel algorithms to converge to a solution using a convergence tolerance of 10^{-12} , when different numbers of processors are being used

5.b Lexicographic Gauss-Seidel on Multiple Processors

Table 6 shows how many iterations it took for Jacobi and Lexicographic Gauss-Seidel to converge with $N = 241$ and a convergence tolerance of 10^{-12} .

There is no change in the number of iterations required by Jacobi to converge from when it was run on a single processor. This is unsurprising, since adding more processors does not actually impact the computation in anyway.

The number of iterations required by Lexicographic Gauss-Seidel is still roughly half the number of iterations it takes for Jacobi to converge, however this property starts to weaken as you increase the number of processors.

Table 6 shows that the number of iterations for Gauss-Seidel increases steadily. I believe the reason for this is the order in which computations and communication is performed on each processor. The *boundary* values each process contains depend on newly computed values from other processors. Since those new values (used for Gauss-Seidel iteration) are not received by a process until *after* the whole iteration, the current iteration uses old values where new values should have been used, which slows the convergence.

In other words, as increase the number of processors, the more and more old values (from iteration k , not $k + 1$) are used for the computation. This causes the Gauss-Seidel algorithm to become closer and closer to Jacobi. If the number of processors was $(N - 1)$, then lexicographic Gauss-Seidel would essentially be Jacobi iteration, because each process would try and use its "new" values, but would actually be using values from the previous iteration k (since the new values for iteration $k + 1$ have not been received from the other processors yet).

# Processors	Convergence Tolerance	# Iterations
1	10^{-12}	114365
1	10^{-10}	87265
1	10^{-8}	60165
1	10^{-6}	33065
4	10^{-12}	114365
4	10^{-10}	87265
4	10^{-8}	60165
4	10^{-6}	33065
7	10^{-12}	114741
9	10^{-12}	116152
11	10^{-12}	114403
7	10^{-6}	33159
9	10^{-6}	33508
11	10^{-6}	33077

Table 7: How many iterations red-black Gauss-Seidel iteration requires with different convergence tolerances

6 Red-Black Gauss-Seidel Iteration

`redblack_gauss.c` contains a parallel implementation of red-black Gauss-Seidel iteration. Table 7 shows how many iterations it takes this algorithm to converge with different convergence tolerances, both on one and four processors.

For one processor, red-black Gauss-Seidel iteration uses the same amount of iterations is the same as lexicographic Gauss-Seidel. Unlike the lexicographic algorithm however, the number of iterations red-black version uses does not increase as you increase the number of processors (as long as the number of processors p makes $\frac{(N-1)}{p}$ an even integer). In other words, red-black Gauss-Seidel scales to more processors and parallelises much better than lexicographic Gauss-Seidel.

Table 7 also contains iteration counts for when I ran the algorithm on a p processors where $\frac{(N-1)}{p}$ is not an even integer. In other words, when the work is not be evenly distributed among the processes (one will have less work than the others). The numbers I used were 7, 9 and 11.

The key difference here is that the number of iterations required to converge *increases slightly*, in a similar fashion to lexicographic Gauss-Seidel. This is due to the fact that there will be some processes where the first node on the first row is actually a *black* node, but the processes itself will treat it as a red node (because it's the first node of the array). This causes a mismatch on the latest values at the boundaries of some processes (values from iteration k used where values from $k + 1$ should be) , which affect the speed at which the algorithm converges.

7 Successive Over-Relaxation

`sor.c` is a modification of `redblack_gauss.c` which makes use of Successive Over-Relation. The relaxation parameter ω which produced the optimal number of iterations for the problem $u = 1, N = 241, Tol = 10^{-12}$ was found.

Table 8 shows the iteration counts for different values of ω . Note that for some values ω , particular very high values, SOR never converged and continuously fluctuated between two different values. These are marked appropriately in the table. $\omega = 1.98$ produced the lowest number of iterations (1339).

w	# Iterations
0.9	138338
1.0	114365
1.1	94536
1.2	77831
1.3	63538
1.4	51143
1.5	40266
1.6	30615
1.7	21959
1.8	14102
1.9	6835
1.91	6128
1.92	5420
1.93	4711
1.94	3995
1.95	3267
1.96	2509
1.97	1653
1.98	1339
1.99	2709
2.0	> 150000 (does not converge)

Table 8: How relaxation parameter affects number of iterations when using a convergence tolerance of 10^{-12} on **four processors**

8 Comparison of the Three Iterative Methods

8.a Comparison on a Single Processor

The runtimes of Jacobi, red-black Gauss-Seidel and red-black Successive Over-Relaxation when solving $u = 1, N = 241$ with a convergence tolerance of 10^{-12} have been recorded in Table 9.

On a single processor, Jacobi and red-black Gauss-Seidel iteration have similar runtimes, with red-black Gauss-Seidel taking a little longer. While red-black Gauss-Seidel took half the iterations as Jacobi, each iteration takes more than twice as long as an one iteration of Jacobi (as shown in Table 9). Therefore, there is no gain on speed with a single processor. Red-black Gauss-Seidel's iterations take longer than Jacobi's because there are two lots of communication per iteration and two passes through the array to update the red and black nodes separately.

A single iteration of red-black SOR takes a similar amount of time as red-black Gauss-Seidel. However, it runs significantly faster overall due to the massive reduction in iterations. On a single processor, red-black SOR is by far the fastest, being over 30 times faster than the other two algorithms.

8.b Comparison on Multiple Processors

On multiple processors, red-black Gauss-Seidel now outperforms Jacobi. This is because it now uses a low enough number of iterations that the extra cost per iteration (compared to Jacobi) does not outweigh the savings in time from less iterations. SOR still beats red-black Gauss-Seidel by a large margin, requiring less iterations and thus, less time, to converge.

By parallelising, the runtimes of each individual iteration has been reduced, but the *relationship* between iteration runtimes for each algorithm remains the same. One iteration of red-black Gauss-Seidel still takes roughly double the time of a single iteration of Jacobi and red-black SOR still has a similar iteration runtime as red-black Gauss-Seidel.

Algorithm	# Processors	Runtime (in seconds)	# Iterations	Runtime per Iteration
Jacobi	1	34.2217	220567	0.0001551533
Red-Black Gauss Seidel	1	43.2007	114365	0.00037774406
Red-Black SOR ($\omega = 1.98$)	1	1.07561	1339	0.0008032935
Jacobi	4	18.0577	220567	0.00008186945
Red-Black Gauss Seidel	4	12.2005	114365	0.00010668036
Red-Black SOR ($\omega = 1.98$)	4	0.334205	1339	0.00024959297

Table 9: Runtime of different algorithms with varying number of processors for $u = 1$, $N = 241$ and a convergence tolerance of 10^{-12}

8.c Shortest Runtime Achieved

The shortest run-time I achieved with a convergence tolerance of 10^{-12} is 0.334205 seconds. This was from using red-black Successive Over-Relaxation (with $\omega = 1.98$ on four processors). This is approximately $\frac{34.22217}{0.334205} \approx 102$ times faster than running Jacobi iteration on a single processor.

NOTE: If the number of processes used to execute SOR is increased to five processors, then the communications overhead of having more processors starts to dominate, making the overall runtime *slower* than four processors.

9 Appendix

9.a Task 1 Code Modifications

The original source code contains a function called `exact()`, which returns the exact solution at point (x, y) . To change the exact solution to $e^{3y} \cos 3x$, the following code was used:

```

1 double exact( double x, double y )
2 {
3     double solution;
4
5     solution = exp(3 * y) * cos(3 * x);
6
7     return (solution);
8 }

```

The convergence tolerance was set to 10^{-12} by changing the `Tol` macro like so:

```
#define Tol 0.000000000001
```

9.b Task 3 Code Modifications

To fix the number of iterations for both `heat_rows.c` and `heat_blocks.c`, I defined a new constant:

```
1 #define MAX_ITER 8000
```

Then I changed the stopping criteria by changing the condition of the algorithm's while-loop:

```
1 while ( iter < MAX_ITER )
```

9.c Task 4 MPD Ring

An MPD ring with the following machines were used to obtain runtimes for tasks 4 and 4.

- csln011
- csln012
- csln013
- csln014

Four processes (i.e. four CPUs) were used for each machine, allowing for a maximum of 16 simultaneous processes.

9.d Task 5 Code Modifications

To implement lexicographic Gauss-Seidel, I simply modified the equation inside the **iteration()** function to use the newly updated values for the left and top neighbouring nodes:

```
1 new[i][j] = 0.25 * ( old[i+1][j] + new[i-1][j] + old[i][j+1] + new[i][j-1] );
```

9.e Task 6 Full Source Code – redblack_gauss.c

```
1 /*
2  * Parallel implementation of 2D Laplace equation solver
3  *   — Jacobi iteration
4  *   — partitioning done by strips of rows
5  *   — non-constant boundary conditions specified
6  *   — exact solution compared with
7  *
8  * Code supplied by Peter Jimack and modified by Matthew Hubbard */
9
10 #include <stdlib.h>
11 #include <stdio.h>
12 #include <math.h>
13 #include <mpi.h>
14
15 /* Define parameter values */
16 #define N 241
17 #define Tol 0.000000000001
18 // 10^-5 = 0.00001
19 // 10^-6 = 0.000001
20 // 10^-8 = 0.00000001
21 // 10^-10 = 0.0000000001
22 // 10^-12 = 0.000000000001
23
24 /* Define the geometric bounds of the domain */
25 #define xmin -2.0
26 #define xmax 2.0
27 #define ymin -2.0
28 #define ymax 2.0
29
```

```

31 enum NodeType {
    NODE_TYPE_RED = 0,
    NODE_TYPE_BLACK
33 };

35 /* Allocate memory for a two-dimensional array (pointers to pointers) */
37 double **matrix( int m, int n )
{
39     int i;
    double **ptr;

41     ptr = (double **) calloc( m, sizeof(double *) );
43     for ( i=0; i<m; i++ )
        ptr[i] = (double *) calloc( n, sizeof(double) );

45     return (ptr);
47 }

49

51 /* Carry out a single Jacobi iteration on the specified range of mesh nodes */
double iteration( double **mesh, int first_row, int last_row, int first_col, int
    last_col, enum NodeType nodeType )
53 {
    double diff, maxdiff=0.0;
55     int i, j;
    int isEven=0;
57     double old=0.0;

59     for ( i=first_row; i<=last_row; i++ )
        for ( j=first_col; j<=last_col; j++ ) {
61         old = mesh[i][j];
        /* The red-black Gauss-Seidel update for node (i,j) */
63         isEven = (i + j) % 2;
        if ((isEven && nodeType == NODE_TYPE_RED) ||
65             (!isEven && nodeType == NODE_TYPE_BLACK)) {
            mesh[i][j] = 0.25 * ( mesh[i+1][j] + mesh[i-1][j] + mesh[i][j+1] + mesh[i][j
67             -1] );
        }

69         /* Update the max norm of  $x^{(k+1)} - x^{(k)}$  */
        diff = mesh[i][j] - old;
71         if ( diff<0 )
            diff = -diff;
73         if ( maxdiff<diff )
            maxdiff = diff;
75     }

77     /* Return the max norm of  $x^{(k+1)} - x^{(k)}$  for the specified range of nodes */
    return (maxdiff);
79 }

81

83 /* Calculate the exact solution at point (x,y) */
double exact( double x, double y )
85 {
    double solution;

87     solution = 1.0;

89     return (solution);
91 }

```

```

93
95 /* Calculate the max norm error in the final approximation on the specified range of
   * mesh nodes
   */
97 double final_error( double **mesh, int start_row, int nrows, int start_col, int ncols,
   double dx, double dy )
{
99     double x, y, diff, maxerr=0.0;
   int i, j;
101
   for ( i=1; i<=nrows; i++ ) {
103       for ( j=1; j<=ncols; j++ ) {
           /* Calculate the geometric coordinates of point (i,j) */
105         x = xmin + dx * (double) (start_row+i-1);
           y = ymin + dy * (double) (start_col+j-1);
107
           /* Update the max norm of approximate - exact */
109         diff = mesh[i][j] - exact( x, y );
           if ( diff<0 )
111             diff = -diff;
           if ( maxerr<diff )
113             maxerr = diff;
115       }
   }
117 /* Return the max norm of approximate - exact for the specified range of nodes */
   return (maxerr);
119 }
121
123 /* Write the approximate solution out to a single data file from process 0
   * — the format is row-by-row */
125 void write_file( int nid, double **mesh, int rowmin, int rowmax, int colmin, int
   colmax, int first_row, int first_col, int nprocs )
{
127     double **full;
   double *send_buffer, *recv_buffer;
129     int row, col, counter, i, j, ip;
   int imin, imax, jmin, jmax;
131     int send_indices[4], recv_indices[4];
   int send_size, recv_size;
133     FILE *fp;
   MPI_Status status;
135
137     /* Fix local row and column indices */
   imin = rowmin - first_row + 1;
139     imax = rowmax - first_row + 1;
   jmin = colmin - first_col + 1;
141     jmax = colmax - first_col + 1;
143
   if ( nid==0 ) {
145       /* Allocate memory on process 0 to store the full solution */
       full = matrix( N+1, N+1 );
147
       /* Insert the solution values calculated on process 0 in to the full array */
149       for ( row=rowmin, i=imin; row<rowmax+1; row++, i++ )
           for ( col=colmin, j=jmin; col<colmax+1; col++, j++ )
151             full[row][col] = mesh[i][j];
153
       /* Receive solution values from the other processes and enter them in to the full
       array */

```



```

155 for ( ip=1; ip<noprocs; ip++ ) {
156     /* Receive the indices indicating the solution values calculated on process ip */
157     MPI_Recv( &recv_indices[0], 4, MPI_INT, ip, 25, MPLCOMM_WORLD, &status );
158
159     rowmin = recv_indices[0]; rowmax = recv_indices[1];
160     colmin = recv_indices[2]; colmax = recv_indices[3];
161
162     /* Calculate the number of solution values to be sent by process ip */
163     recv_size = (rowmax-rowmin+1)*(colmax-colmin+1);
164     recv_buffer = (double *) calloc( recv_size, sizeof(double) );
165
166     /* Receive the solution values calculated on process ip */
167     MPI_Recv( &recv_buffer[0], recv_size, MPLDOUBLE, ip, 15, MPLCOMM_WORLD, &
168     status );
169
170     /* Place the solution values in the message in to the full array, in order */
171     counter = 0;
172     for ( row=rowmin; row<rowmax+1; row++ )
173         for ( col=colmin; col<colmax+1; col++ ) {
174             full[row][col] = recv_buffer[counter];
175             counter++;
176         }
177     free( recv_buffer );
178 }
179
180 /* Write out the full array to heat_solution.dat row by row */
181 fp = fopen( "heat_solution.dat", "wt" );
182
183 for ( row=0; row<N+1; row++ )
184     for ( col=0; col<N+1; col++ )
185         fprintf( fp, "%6d %6d %8.6f\n", row, col, full[row][col] );
186
187 fclose( fp );
188
189 free( full );
190 }
191 else {
192     send_indices[0] = rowmin; send_indices[1] = rowmax;
193     send_indices[2] = colmin; send_indices[3] = colmax;
194
195     /* Send the indices indicating the solution values to be sent to process 0 */
196     MPI_Send( &send_indices[0], 4, MPI_INT, 0, 25, MPLCOMM_WORLD );
197
198     /* Calculate the number of solution values to be sent to process 0 */
199     send_size = (imax-imin+1)*(jmax-jmin+1);
200     send_buffer = (double *) calloc( send_size, sizeof(double) );
201
202     /* Place the solution values to be sent in to the message in a specific order */
203     counter = 0;
204     for ( i=imin; i<imax+1; i++ )
205         for ( j=jmin; j<jmax+1; j++ ) {
206             send_buffer[counter] = mesh[i][j];
207             counter++;
208         }
209
210     /* Send the solution values calculated on this process to process 0 */
211     MPI_Send( &send_buffer[0], send_size, MPLDOUBLE, 0, 15, MPLCOMM_WORLD );
212
213     free( send_buffer );
214 }
215 }

```

```

void synchronise_nodes(double** nodeValues, int nrows, int ncols, int nid, int nprocs
)
217 {
218     MPI_Status status;
219     if ( nid<nprocs-1 ) {
220         MPI_Send( &nodeValues[nrows][1], ncols, MPLDOUBLE, nid+1, 10, MPLCOMM_WORLD
);
221         MPI_Recv( &nodeValues[nrows+1][1], ncols, MPLDOUBLE, nid+1, 20, MPLCOMM_WORLD,
&status );
222     }
223     if ( nid>0 ) {
224         MPI_Send( &nodeValues[1][1], ncols, MPLDOUBLE, nid-1, 20, MPLCOMM_WORLD );
225         MPI_Recv( &nodeValues[0][1], ncols, MPLDOUBLE, nid-1, 10, MPLCOMM_WORLD, &
status );
226     }
227 }

229
/* Main program */
231 int main( int argc, char **argv )
{
232     double **mesh, **tmp;
233     double diff, MaxDiff, MaxDiffG, MaxErr, MaxErrG;
234     double dx, dy, x, y;
235     double start_time=0.0, end_time=0.0;
236     int first_row, nrows, first_col, ncols;
237     int nprocs, nid, remainder, i, j, iter;
238     int rowmin, rowmax, colmin, colmax;
239     MPI_Status status;
240     MPI_Request req_send10, req_send20, req_recv10, req_recv20;

241
242
/* Initialise for MPI */
243 MPI_Init( &argc, &argv );
244 MPI_Comm_rank( MPLCOMM_WORLD, &nid );
245 MPI_Comm_size( MPLCOMM_WORLD, &nprocs );

246
247
/* Calculate the mesh size */
248 dx = ( xmax - xmin ) / (double) N;
249 dy = ( ymax - ymin ) / (double) N;

250
251
/* Calculate the first row and the number of rows allocated to this process */
252 first_row = 1;
253 remainder = (N-1)%nprocs;
254 nrows = (N-1-remainder)/nprocs;
255 /* Care is taken with any left over rows after dividing the number of rows by the
number of
* processes: they are allocated one-by-one to the first few processes to avoid
having one
* process taking much more work than all of the others
*/
256 for ( i=0; i<nid; i++ ) {
257     if ( i<remainder )
258         first_row = first_row + nrows + 1;
259     else
260         first_row = first_row + nrows;
261 }
262 if ( nid<remainder )
263     nrows++;

264
265
/* Calculate the first column and the number of columns allocated to this process */
266 first_col = 1;
267 ncols = N-1;

```

```

275  /* Allocate memory for the local work arrays */
277  mesh = matrix( nrows+2, ncols+2 );

279
281  if ( nid==0 )
      start_time = MPI_Wtime();

283
285  /* Apply boundary conditions at the ymin and ymax boundaries */
287  for ( i=0; i<=nrows+1; i++ ) {
289      x = xmin + dx * (double) (first_row+i-1);
291      y = ymin;
      mesh[i][0] = exact( x, y );
      y = ymax;
      mesh[i][ncols+1] = exact( x, y );
  }

293  /* Apply boundary conditions at the xmin and xmax boundaries */
295  if ( nid==0 )
      for ( j=1; j<=ncols; j++ ) {
297          x = xmin;
299          y = ymin + dy * (double) (first_col+j-1);
          mesh[0][j] = exact( x, y );
      }
301  if ( nid==nprocs-1 )
      for ( j=1; j<=ncols; j++ ) {
303          x = xmax;
          y = ymin + dy * (double) (first_col+j-1);
          mesh[nrows+1][j] = exact( x, y );
305      }

307  // Just set so one iteration of the loop is guaranteed to fire
MaxDiffG = Tol * 2;
309  iter = 1;
while ( MaxDiffG>Tol ) { /* The error estimate is still too large */
311      MaxDiff = iteration(mesh, 1, nrows, 1, N-1, NODETYPE.RED);
      synchronise_nodes(mesh, nrows, ncols, nid, nprocs);
313      diff = iteration(mesh, 1, nrows, 1, N-1, NODETYPE.BLACK);
      synchronise_nodes(mesh, nrows, ncols, nid, nprocs);
315      if ( diff > MaxDiff )
          MaxDiff = diff;

317
319      /* Estimate the error for the complete update */
      MPI_Allreduce( &MaxDiff, &MaxDiffG, 1, MPLDOUBLE, MPLMAX, MPLCOMM.WORLD );

321      /* Write out the error estimate every 1000 iterations */
      iter++;
323      if ( nid==0 && iter%1000==0 )
          fprintf( stdout, "Iteration = %i, MaxDiffG = %12.5e\n", iter, MaxDiffG );
325  }

327  /* Compute the error in the final approximation */
MaxErr = final_error( mesh, first_row, nrows, first_col, ncols, dx, dy );
329  MPI_Allreduce( &MaxErr, &MaxErrG, 1, MPLDOUBLE, MPLMAX, MPLCOMM.WORLD );

331  if ( nid==0 ) {
      fprintf( stdout, "Iteration = %i, MaxDiffG = %12.5e\n", iter, MaxDiffG );
333      fprintf( stdout, "Error = %12.5e\n", MaxErrG );

335      end_time = MPI_Wtime();
      fprintf( stdout, "Execution time in seconds = %12.5e\n", end_time-start_time );
337  }

```

```

339  /* Assign bounds for the local and global indices of the rows and columns to
340     indicate
341     * the nodes of the mesh that are updated by this process, purely for the purpose of
342     * ending all of the solution to process 0 for writing to a single file
343     */
344  rowmin = first_row;
345  rowmax = first_row + nrows - 1;
346  colmin = first_col;
347  colmax = first_col + ncols - 1;
348
349  /* The data values to be written to file should also include the boundary values, so
350     * adjust the row and column bounds on the appropriate processors to include these
351     */
352  if ( nid==0 ) {
353      rowmin--;
354  }
355  if ( nid==nprocs-1 ) {
356      rowmax++;
357  }
358  colmin--;
359  colmax++;
360
361  /* Write the data to file */
362  write_file( nid, mesh, rowmin, rowmax, colmin, colmax, first_row, first_col, nprocs
363             );
364
365  /* Free allocated memory */
366  free( mesh );
367
368  /* Finish up */
369  MPI_Finalize();
370
371  return (0);
372 }

```